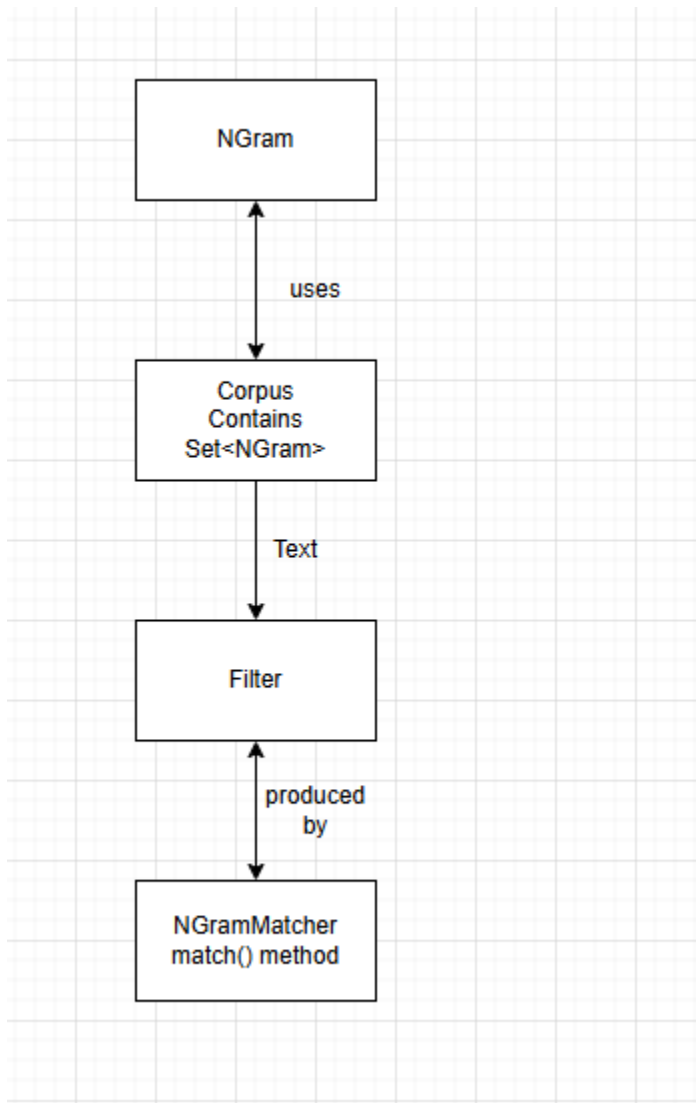


General Idea

The system is designed as a modular, object-oriented Java project.

- NGram is an immutable value object representing a word.
- Corpus is a container of valid NGram words; its internal collection is immutable.
- Filter encapsulates the feedback rule (similar to Wordle's green/yellow/gray) as a Predicate over NGrams.
- NGramMatcher implements a three-pass matching algorithm, returning a Filter that specifies which letters in a guess are exact, misplaced, or absent.
- MatchleScorer uses the matcher and corpus to score guesses (worst-case and average-case) and select optimal guesses.



2. Detailed Class Descriptions

2.1 NGram

- Attributes:
 - private final ArrayList<Character> ngram;
Stores the characters of the word in order.
 - private final Set<Character> charset;
A helper set containing unique letters (for quick membership testing).
- Routines:
 - static NGram from(String str): NGram
Input: A non-null String.
Precondition: str \neq null (throws NullPointerException if violated).
Postcondition: Returns a new immutable NGram with size equal to string length.
 - static NGram from(List<Character> chars): NGram
Input: A non-null List of Characters containing no null elements.
Precondition: List and all elements are non-null.
Postcondition: Returns a new NGram instance.
 - Character get(int index): Character
Input: An int index.
Precondition: $0 \leq \text{index} < \text{ngram.size}()$ (throws IndexOutOfBoundsException if violated).
Postcondition: Returns the character at the given index.
 - int size(): int
Returns the number of characters in this NGram.
 - boolean matches(IndexedCharacter c): boolean
Input: An IndexedCharacter (record with index and character).
Precondition: c \neq null.
Postcondition: Returns true if the character at c.index equals c.character.
 - boolean contains(char c): boolean
Returns true if the character c is contained in the NGram.
 - boolean containsElsewhere(IndexedCharacter c): boolean
Input: An IndexedCharacter.
Precondition: c \neq null.

Postcondition: Returns true if c.character is present in a different index than c.index.

- Iterator and Stream methods:
Provide iteration over IndexedCharacter objects.
 - Error Handling:
 - All public routines check for null inputs and enforce index bounds.
 - Defensive copying is used in factory methods to preserve immutability.
 - Testing:
 - Unit tests should verify correct construction, equality, immutability, iterator behavior, exception cases (null/invalid index), and correct responses from methods like matches, contains, containsElsewhere.
-

2.2 Corpus

- Attributes:
 - private final Set<NGram> corpus;
An unmodifiable set of NGram objects (built by copying the input set).
 - private final int wordSize;
All words in the corpus have this common length.
- Routines:
 - Set<NGram> corpus(): Set<NGram>
Returns the unmodifiable set of NGrams.
 - int wordSize(): int
Returns the common word length.
 - int size(): int
Returns the number of NGrams in the corpus.
 - boolean contains(NGram ngram): boolean
Input: Non-null NGram.
Precondition: ngram ≠ null.
Returns whether the corpus contains the NGram.
 - Iterator<NGram> iterator() and Stream<NGram> stream()
Return iterators or streams over the corpus.

- long size(Filter filter): long
Input: A non-null Filter.
Returns the count of NGrams in this corpus satisfying the filter.
 - Nested Builder Class:
 - Attributes:
 - private final Set<NGram> ngrams;
A modifiable set used to accumulate NGrams before building the final Corpus.
 - Routines:
 - Builder add(NGram ngram): Builder
Input: Non-null NGram; adds it to the builder.
 - Builder addAll(Collection<NGram> collection): Builder
Adds all non-null NGrams from the collection.
 - boolean isConsistent(Integer wordSize): boolean
Input: Non-null Integer.
Returns true if all NGrams in the builder have the given wordSize.
 - Corpus build(): Corpus
Precondition: All NGrams are non-null and of equal length.
Postcondition: Returns an immutable Corpus or throws IllegalStateException if inconsistent.
 - Error Handling:
 - All public methods check for null.
 - The Builder's build() method enforces consistent state.
 - Testing:
 - Unit tests should create corpora using the Builder.
 - Test immutability (try to modify the corpus, expect an exception).
 - Verify proper behavior when null values are added or inconsistent sizes appear.
-

2.3 Filter

- Attributes:

- private final Predicate<NGram> predicate;
The internal condition that a candidate NGram must meet.
 - Routines:
 - static Filter from(Predicate<NGram> predicate): Filter
Input: A non-null Predicate.
Precondition: predicate ≠ null.
Returns: A Filter encapsulating the predicate.
 - boolean test(NGram ngram): boolean
Input: A non-null NGram.
Returns: The result of applying the predicate.
 - Filter and(Optional<Filter> other): Filter and Filter and(Filter other): Filter
Combine this Filter with another using logical AND.
 - Error Handling:
 - Every method validates its input.
 - The internal predicate is set only if not null.
 - Testing:
 - Unit tests should check that filters behave as intended (true/false for given NGrams), including combining filters (and) and handling null parameters (which should throw exceptions).
-

2.4 NGramMatcher

- Attributes:
 - private final NGram key; – The secret word to be matched.
 - private final NGram guess; – The player's guess.
- Routines:
 - static NGramMatcher of(NGram key, NGram guess): NGramMatcher
Input: Two non-null NGrams.
Precondition: key and guess ≠ null.
Returns: A new NGramMatcher instance.
 - Filter match(): Filter
Returns: A Filter representing the feedback comparing key and guess. *Algorithm:* (See pseudocode below)

- *Precondition:* `key.size() == guess.size()`; if not, returns `Filter.FALSE`.
- *Postcondition:* The returned `Filter` combines conditions on each letter position:
 - **Exact Matches (Green):** For every `i` where `key[i]` equals `guess[i]`, add condition that candidate must have that letter at position `i`.
 - **Misplaced Matches (Yellow):** For each unmatched guess letter, search for a matching letter in an unmatched position of `key`; add condition that candidate should contain the letter elsewhere.
 - **Absent (Gray):** For guess letters that remain unmatched, add condition that candidate must not contain that letter.
- **Helper Methods (private):**
 - `doExactMatches(...)`
For each index `i`, if `key[i]` equals `guess[i]`, mark as matched and add an exact match `Filter`.
 - `doMisplacedMatches(...)`
For each index not matched, call `handleMisplacedAtIndex(...)` to check for misplaced letters, and add a misplaced `Filter`.
 - `doAbsentCharacters(...)`
For each unmatched position, add an absent `Filter`.
 - `handleMisplacedAtIndex(int i, int n, boolean[] keyMatched, boolean[] guessMatched, List<Filter> partialFilters)`
Processes a single guess index `i` to find a match in any unmatched position `j` in `key`; if found, marks them as matched and adds a misplaced filter.
- **Error Handling:**
 - The factory method uses `Objects.requireNonNull` for `key` and `guess`.
 - Early exit in `match()` when lengths differ.
- **Testing:**
 - Unit tests should cover:
 - Different lengths (expect `Filter.FALSE`).
 - Exact match only.
 - Misplaced match (with duplicate letters).

- Absent letters.
- Private helper methods can be indirectly tested via the public match() method.
- Test invalid cases (null inputs; out-of-bound indices through NGram methods)

Pseudocode Summary for NGramMatcher.match():

function match(guess, target) -> Filter:

if guess.length != target.length:

return Filter.FALSE

initialize keyMatched[] and guessMatched[] of size n to false

initialize partialFilters as empty list

// PASS 1: Exact matches

for each index i in 0 to n-1:

if guess[i] == target[i]:

mark keyMatched[i] and guessMatched[i] true

add exact match filter at i to partialFilters

// PASS 2: Misplaced matches

for each index i in 0 to n-1:

if guessMatched[i] is false:

call handleMisplacedAtIndex(i)

for each j in 0 to n-1:

if keyMatched[j] is false and target[j] equals guess[i]:

mark keyMatched[j] and guessMatched[i] true

add misplaced filter for i to partialFilters

break

// PASS 3: Absent letters

for each index i in 0 to n-1:

if guessMatched[i] is false:

add absent filter for guess[i] to partialFilters

Combine all partialFilters using logical AND

return the final Filter

2.5 MatchleScorer

- Attributes:
 - private final Corpus corpus; – The immutable dictionary of NGrams.
- Routines:
 - MatchleScorer(Corpus corpus): void
Input: A non-null, non-empty Corpus.
Precondition: corpus \neq null and corpus.size() > 0.
Postcondition: Stores the corpus in an immutable field.
 - long score(NGram key, NGram guess): long
Inputs: Non-null key and guess.
Precondition: key and guess are not null.
Process: Uses NGramMatcher.of(key, guess).match() to generate a Filter.
Return: The count of NGrams in the corpus that satisfy the feedback filter.
 - long scoreWorstCase(NGram guess): long
Input: A non-null guess.
Return: The maximum score over all keys in the corpus (worst-case scenario).
 - long scoreAverageCase(NGram guess): long
Input: A non-null guess.
Return: The sum of scores over all keys in the corpus.
 - NGram bestGuess(ToLongFunction<NGram> criterion): NGram
Input: A non-null criterion function mapping a candidate guess to a long score.
Return: The NGram in the corpus with the lowest score by the criterion.
 - NGram bestWorstCaseGuess(): NGram
Delegates: Calls bestGuess(this::scoreWorstCase).

- NGram `bestAverageCaseGuess()`: NGram
Delegates: Calls `bestGuess(this::scoreAverageCase)`.
 - Error Handling:
 - All methods use defensive checks (`Objects.requireNonNull`) for inputs.
 - The constructor throws an exception if corpus is null or empty.
 - Testing:
 - Unit tests for all scoring methods verifying expected outputs.
 - Tests verifying that passing null causes `NullPointerException`.
-

3. Error Handling and Defensive Strategies

- Defensive Checks:
 - Every public method uses `Objects.requireNonNull(...)` to check for nulls.
 - Preconditions on methods (such as matching lengths in `NGramMatcher.match`) are checked early; if violated, a known (false) value or exception is returned.
 - Immutability:
 - NGram and Corpus objects are immutable.
 - Defensive copies are made in constructors (e.g., `Corpus` and `NGram`).
 - Clear Exceptions:
 - Methods throw specific exceptions (`NullPointerException`, `IllegalStateException`, `IndexOutOfBoundsException`) with descriptive messages.
-

4. Testing Strategy

- Unit Tests:
 - Write a dedicated test class for each core component:
 - `NGramTest`: Verify construction (from `String/List`), correct element access, iteration, `equals/hashCode`, boundary conditions.
 - `CorpusTest`: Verify that the Builder correctly accumulates NGrams, enforces consistency, and that public methods return the expected unmodifiable view.

- FilterTest: Verify that the Filter factory methods, test(), and the logical AND methods behave as expected and enforce null checks.
- NGramMatcherTest: Verify that match() correctly computes feedback (exact, misplaced, absent) for various key-guess pairs.
- MatchleScorerTest: Verify scoring methods (score, worst-case, average-case, best guess methods) return expected values; also test defensive behavior against null inputs.
- Integration Tests:
 - Simulate a game loop where a series of guesses narrow down the Corpus until the secret is found.
 - Ensure that filtering and scoring are consistent across rounds.
- Stress Test:
 - Create a large Corpus (e.g., using thousands of 5-letter words from an extended dictionary).
 - Run through many simulated games in a loop, ensuring that:
 - Performance remains acceptable.
 - No memory leaks occur (by monitoring memory usage over iterations).
 - All defensive conditions remain true, and no invalid state is reached.
 - This single stress test would run a loop (perhaps tens of thousands of iterations) that randomly selects a secret, simulates guesses using the scoring functions, and verifies that the system consistently converges on a valid answer.
- Test Hooks:
 - For private methods (especially in NGramMatcher), design your tests to cover them indirectly via the public match() method.
 - If needed, use reflection (or package-private visibility) to test edge cases for private methods, though ideally the public behavior covers the important logic.

5. Single Stress Test Description

Stress Test Outline:

- Purpose: Test the performance and stability of the scoring and matching system with a large corpus and many iterations.

- Setup:
 - Load a large dictionary of NGrams (e.g., 10,000 words of a fixed length).
 - Build a Corpus from this dictionary.
- Test Process:
 - For a fixed number of iterations (e.g., 10,000 rounds), randomly select a secret key from the Corpus.
 - For each iteration, simulate a guess sequence:
 1. Use MatchleScorer to compute worst-case or average-case scores for a set of candidate guesses.
 2. Select the best guess.
 3. Use NGramMatcher.match() to generate a feedback Filter.
 4. Filter the Corpus and update candidate list.
 - Verify that after a sequence of rounds (or simulation), the candidate set is reduced to one word—the chosen secret.
 - Record timing and monitor memory usage.
- Validation:
 - Ensure each iteration completes without errors or exceptions.
 - Confirm that the system's performance scales linearly (or better) with the number of words.
 - The stress test should run unattended, logging summary metrics (total runtime, average time per round, memory usage).